MUSAN: A Music, Speech, and Noise Corpus

David Snyder[1,*], Guoguo Chen[1], and Daniel Povey[1]

[1]Center for Language and Speech Processing, The Johns Hopkins University, Baltimore, MD 21218, USA
[*]david.ryan.snyder@gmail.com

ABSTRACT

This report introduces a new corpus of music, speech, and noise. This dataset is suitable for training models for voice activity detection (VAD) and music/speech discrimination. Our corpus is released under a flexible Creative Commons license. The dataset consists of music from several genres, speech from twelve languages, and a wide assortment of technical and non-technical noises. We demonstrate use of this corpus for music/speech discrimination on Broadcast news and VAD for speaker identification.

1 Introduction

Classification of audio into speech, music, or other nonspeech categories has several practical applications. Voice activity detection (VAD) is often an essential preprocessing step for other speech technologies, such as speech recognition, speaker diarization, or speaker verification. In some applications, on hold music needs to be detected and removed.

Gaussian mixture models (GMM) are a simple and popular model for VAD and music/speech discrimination [2, 7]. In [2], several frame-level VAD techniques were compared for speaker verification, and an unsupervised GMM-based VAD was found to produce the best results on NIST SRE 2008. We also make use of GMM-based systems to demonstrate this corpus. Moreover, our focus is on providing the data; therefore, we do not explore complex models or sophisticated features.

Most publicly available corpra for music and speech discrimination that provide raw audio do not address the copyright of the data sources nor appear to have permission to redistribute the data. For instance, the GTZAN Music/Speech dataset [8] is widely used, but it appears that permission was not given by the copyright holders to redistribute the work. Other corpra, such as the Million Song database [3] circumvent intellectual property issues by providing only features. However, this limits the user to building systems based only on the provided feature type. In contrast, our corpus is compiled from Creative Commons and US Public Domain sources, so we are free to redistribute the original audio.

2 The corpus

The corpus consists of approximately 109 hours of audio that is in the US Public Domain or under a Creative Commons license. The directories are partitioned into speech, music, and noise and by data source (i.e., the website we downloaded the content from). It is freely available at OpenSLR. All audio is formatted as 16kHz WAV files.

Each subdirectory contains a LICENSE file which connects a WAV file to its respective license and provides attribution appropriate to the license type. For example, a music entry in a LICENSE file may contain the filename, title, artist, a url to the source, and a summary of the license. All files in this corpus fall under a Creative Commons license or are considered to be in the US Public Domain. To broaden the use of this corpus, we only include content that allows for commercial use. Please refer to the Creative Commons licenses for more information.

Most directories contain an ANNOTATIONS file which provides potentially useful metadata. For example, music is annotated for the presence or absence of vocals and by genre(s). The READMEs in each subdirectory describe the types of annotations in more detail.

2.1 Speech

This portion of the corpus consists of about 60 hours of speech. It contains 20 hours and 21 minutes of read speech from Librivox, all of which are in the Public Domain. Each WAV file is an entire chapter of a book, read by one speaker. Approximately half of the Librivox recordings are in English, and the remainder are from eleven other languages. Annotations for speaker and language are provided. The rest of the speech portion consists of 40 hours and 1 minute of US government (federal and various states) hearings, committees and debates that are believed to be in the US Public Domain. These files have been obtained from the Internet Archive and the Missouri Channel senate archives. These recordings are entirely in English.

2.2 Music

The music portion of the corpus has been downloaded from Jamendo, Free Music Archive, Incompetech, and HD Classical Music. It is 42 hours and 31 minutes. The music is divided into Western art music (e.g., Baroque, Romantic, and Classical) and popular genres (e.g., jazz, bluegrass, hiphop, etc). Annotations for genre, artist, and the presence or absence of vocals are provided. For Western art music, the composer is also provided. These files have all been released under some form of the Creative Commons license. Care was taken to ensure that any of these files can be used for commercial purposes.

2.3 Noise

This portion of the corpus contains 929 files of assorted noises, with a total duration of about 6 hours. These range from technical noises, such as DTMF tones, dialtones, fax machine noises, and more, as well as ambient sounds, such as car idling, thunder, wind, footsteps, paper rustling, rain, animal noises, etc. We do not include recordings with intelligible speech. However, some recordings are crowd noises with indistinct voices. The files were downloaded from Free Sound and Sound Bible. The Free Sound part of the corpus is Public Domain material and the Sound Bible part is CC licensed.

3 Experiments

We demonstrate use of this corpus by training several simple systems for music/speech discrimination and voice activity detection. Our experiments use the Kaldi ASR toolkit [6].

3.1 Music/speech discrimination

We train two GMMs on the speech and music portions of the corpus. The features are 20 MFCCs with sliding window mean normalization. The first through fourth order deltas are appended to the MFCCs. We compare with an identical system trained on the GTZAN dataset.

We test the systems on Broadcast news [4]. This evaluation is similar to [7]. The test audio is separated into speech and music segments. Any overlapping segments were excluded from the evaluation. Classification is performed on the entire segment, by taking the majority of the frame-level decisions. We evaluate at the equal error-rate (EER) operating point.

3.2 Music/speech discrimination results

EER(%)/K	4	8	16	32	64	128
MUSAN	4.43	3.85	3.75	3.85	3.95	4.14
GTZAN	3.85	3.95	3.85	4.05	4.05	3.95

Table 1. A comparison of systems trained on the MUSAN and GTZAN corpra. We calculate the EER(%) using GMMs with $K \in \{4, 8, 16, 32, 64, 128\}$ components.

We experiment with GMMs of varying sizes. In Table 1 we see that, overall, performance isn't affected much by the GMM size, and similar performance is achieved with the models trained on the GTZAN and MUSAN corpra.

## 3.3 Speaker recognition

Figure 1. Speaker recognition schema.

The speaker recognition system is a typical i-vector-based system. It uses a GMM universal background model (UBM) and a PLDA backend (See Figure 1). It is based on the GMM system described in [5]. The raw features are 20 MFCCs with a 25ms frame-length. The front end performs a sliding window mean-normalization over a 3 second window, concatenates delta and acceleration, and uses the voice activity detection (VAD) decisions to remove unvoiced frames.

Our evaluation is based on the core NIST SRE 2010 [1] evaluation. To better expose the effects of voice activity detection, we modify the original evaluation by restricting the amount of speech available at test time. In the original evaluation, the utterances are all approximately 5 minutes long. In this test, only the first n seconds of speech are available to the system at test time. This simulates a practical scenario in which recognition needs to be performed quickly at test time, and better exposes the benefits of a more accurate VAD decisions.

3.3.1 Energy VAD
Our baseline system uses a simple energy VAD. The energy VAD classifies a frame as speech or nonspeech by using information about the average log-energy in a given window, centered around the current frame.

3.3.2 GMM and Energy VAD
This system uses a GMM-based VAD trained on this corpus in addition to the energy-based VAD described in Section 3.3.1.

Three full-covariance GMMs, each with 8 components, are trained on the speech, music, and noise portions of the corpus. We use only the music without vocals, and only the Librivox portion of the speech data. The features are identical to those used for music/speech discrimination in Section 3.1.

To classify a frame a speech or nonspeech, we first apply the energy-based VAD. All of the frames classified as speech are then refined by the GMM-based VAD. We reclassify a frame by selecting the GMM with the highest posterior. The priors are 0.07 for music, 0.75 for speech, and 0.18 for noise, and were selected by tuning on an out-of-domain dataset. Frames classified as music or noise are mapped to nonspeech.

VAD/EER(%)	Max	60s	10s	5s	3s	2s	1s
Energy	2.54	2.99	6.52	9.98	14.17	17.48	23.10
GMM+Energy	2.37	2.65	5.01	8.06	11.23	14.50	20.25
Rel. Improv.	6.70	11.37	23.16	19.24	20.75	17.05	12.34

Table 2. Speaker recognition performance on the core NIST SRE 2010 evaluation with and without the GMM-based VAD.

3.4 VAD Results
In Table 2 we see that the addition of a GMM-based VAD improves performance on the speaker recognition task. The speaker verification results are improved for all amounts of speech, but the benefit is greater when less speech is available.

4 Acknowledgments

5 Conclusion

We introduced a corpus for training music/speech discrimination and voice activity detection (VAD) systems. This corpus was collected entirely from US Public Domain and Creative Commons sources, and provides the raw audio. We showed that this dataset can be used to train simple classifiers for music/speech discrimination and also for frame-level VAD.

References

1. The NIST year 2010 speaker recognition evaluation plan. http://www.itl.nist.gov/iad/mig/tests/sre/2010/, 2010.
2. J. Alam, P. Kenny, P. Ouellet, T. Stafylakis, and P. Dumouchel. Supervised/unsupervised voice activity detectors for text-dependent speaker recognition on the rsr2015 corpus. In *Proc of the Odyssey speaker and language recognition workshop,(Odyssey'14)*, pages 123–30, 2014.
3. T. Bertin-Mahieux, D. P. Ellis, B. Whitman, and P. Lamere. The million song dataset. In *Proceedings of the 12th International Conference on Music Information Retrieval (ISMIR 2011)*, 2011.
4. J. F. W. F. David Graff, John Garofolo and D. Pallett. 1996 english broadcast news speech (hub4). Linguistic Data Consortium, Philadelphia, 1997.

5. D. P. David Snyder, Daniel Garcia-Romero. Time delay deep neural network-based universal background models for speaker recognition. In *ASRU*, 2015.

6. D. Povey, A. Ghoshal, G. Boulianne, L. Burget, O. Glembek, N. Goel, M. Hannemann, P. Motlíček, Y. Qian, P. Schwarz, et al. The Kaldi speech recognition toolkit. In *Proceedings of the Automatic Speech Recognition & Understanding (ASRU) Workshop*, 2011.

7. G. Sell and P. Clark. Music tonality features for speech/music discrimination. In *Acoustics, Speech and Signal Processing (ICASSP), 2014 IEEE International Conference on*, pages 2489–2493. IEEE, 2014.

8. G. Tzanetakis. Gtzan music/speech. http://marsyasweb.appspot.com/download/data_sets/. Accessed: 2015-11-24.